



Lecture 13: Graphical Models & Bayesian Inference Networks

Prof. Krishna R. Pattipati
Dept. of Electrical and Computer Engineering
University of Connecticut
Contact: krishna@engr.uconn.edu (860) 486-2890

Spring 2021
April 30, 2021



Lecture Outline

- Graphical Models
- Bayesian Inference in Graphical Models
- Forward-Backward Methods of Inference
- Advanced Methods
- Summary



Reading List

- Bishop, Chapters 8 and 11
- Murphy, Chapters 19-24
- Theodoridis, Chapter 15



Bayes' Theorem

- Basic Axioms of probability
 - Probability of event A , $P(A) \in [0, 1]$
 - $P(A) = 1 \Leftrightarrow A$ is certain
 - $P(A \cup B) = P(A \text{ or } B) = P(A) + P(B) - P(A \cap B)$
 - Bayes' theorem
 - $P(AB) = P(A|B)P(B) = P(B|A)P(A)$
 - $$P(A|B) = \frac{P(B|A)P(A)}{P(B)}$$
 - Interested in A
 - Begin with a *a priori* probability $P(A)$ for our belief about A
 - Observe B
 - Bayes' theorem provides the revised belief about A , that is, the posterior probability $P(A/B)$




Causality and Inference

- **Likelihood of A:** The quantity $P(B/A)$, as a function of varying A for a fixed B
- $\text{posterior} \propto \text{prior} \times \text{likelihood}$
 $\Rightarrow P(A/B) \propto P(A) \cdot P(B/A)$

We are inferring A given data B

- Graphical representation of the cause-effect process



A causes B (A is the cause and B is the effect)

- Why Graphical Structures?
 - Provide a representation for the joint distribution of a set of variables in terms of conditional and prior probabilities
 - Orientation of the arrows represent influence (causation)
 - Corresponding conditional probabilities are obtained from data or elicited from an expert



Bayesian Inference as Operations on a Graph

- Probabilistic (Bayesian) Inference
 - When data is observed, inferencing is required
 - Involves calculating marginal probabilities of causes conditioned on the observed data using Bayes' theorem
 - Diagrammatically equivalent to reversing one or more of the arrows



“ From the observed effect B to the inferred cause A ”

$$P(AB) = P(B) \cdot P(A/B)$$


- Example 1

$$P(A | B = b) = \frac{P(B = b | A)P(A)}{P(B = b)}$$

BN provide a means to infer the distributions of unobserved variables based on observed ones

Tail-to-Tail Dependency

- Example 2 (a)



“tail-to-tail” dependency

- Naïve Bayes
- iid observations
- prediction (sufficient statistics)

$$P(ABC) = P(A)P(B | A)P(C | A)$$

When A is not observed, B & C are **dependent**.
When A is observed, B & C are **conditionally independent!**

- “Factorization of Joint Distribution”, suppose

- Know $P(A)$, $P(B|A)$ and $P(C|A)$
- Observed $B = b$
- Calculate $P(C|B = b)$

$$B \perp C | A$$

"B is independent of C given A"

why?

$$P(B, C) = \sum_A P(B | A)P(C | A)P(A) \neq P(B)P(C)$$

$$P(B, C | A) = \frac{P(A, B, C)}{P(A)} = P(B | A)P(C | A)$$

- One way of computing it:

1. Calculate $P(ABC)$
2. Compute $P(B) = \sum_A \sum_C P(A, B, C) \Rightarrow$ can get $P(B = b)$
3. Compute $P(B, C) = \sum_A P(A, B, C) \Rightarrow$ can get $P(C, B = b)$
4. Calculate $P(C | B = b) = \frac{P(C, B = b)}{P(B = b)}$

$$P(A, B, C) = \frac{\overset{\text{clique}}{P(A, B)} \overset{\text{clique}}{P(A, C)}}{\underset{\text{separator}}{P(A)}}$$



Exploiting Dependency Structure

- Problem:

Need to compute $|A||B||C|$ entries to compute $P(A,B,C)$

– If $|A|=|B|=|C|=10 \rightarrow$ need 1000 entries

- Alternate way: Exploit the graph structure

1. Calculate $P(A|B=b) = \frac{P(B=b|A)P(A)}{P(B=b)}$ using Bayes' rule,

where $P(B=b) = \sum_A P(B=b|A)P(A) \rightarrow$ arc reversal or inferencing

2. Find
$$\begin{aligned} P(C|B=b) &= \sum_A P(C,A|B=b) \\ &= \sum_A P(C|A,B=b)P(A|B=b) \\ &= \sum_A P(C|A)P(A|B=b) \end{aligned}$$


▪ Advantage: Need to store only 100 entries when $|A|=|B|=|C|=10$

Head-to-Tail Dependency

- Example 2 (b)

“head-to-tail” dependency

- Markov Chains
- HMMs



$$P(ABC) = P(B)P(A|B)P(C|A)$$

$$\text{Also, } P(ABC) = \frac{P(AB)P(CA)}{P(A)}$$

$\{AB\}, \{CA\}$ “cliques”

A “separator”

$P(AB), P(CA)$ are “clique potentials”

$P(A)$ “separator potential”

When A is not observed, B & C are **dependent**.

When A is observed, B & C are **conditionally independent!**

$$B \perp C \mid A$$

“B is independent of C given A”


why?

$$P(B, C) = P(B) \sum_A P(A|B)P(C|A) \neq P(B)P(C)$$

$$P(B, C \mid A) = \frac{P(A, B, C)}{P(A)} = \frac{P(A|B)P(B)P(C|A)}{P(A)} = P(B|A)P(C|A)$$

Head-to-Head Dependency

- Example 2 (c)



“head-to-head” dependency

$$P(ABC) = P(B)P(C)P(A|B, C)$$

When A is not observed, B & C are **independent!!**
When A is observed, B & C are **conditionally dependent!!**

$$B \perp C | \emptyset$$

"B is independent of C given no evidence"

why?

$$P(B, C) = \sum_A P(B)P(C)P(A|B, C) = P(B)P(C)$$

They are not independent given A


$$P(B, C | A) = \frac{P(A, B, C)}{P(A)} = \frac{P(B)P(C)P(A|B, C)}{P(A)}$$

When A is not observed, A blocks the path from B to C. However, when A is observed, it **unblocks** the path from B to C \Rightarrow *they become dependent*


D-Separation

- D (dependency)-separation: ideas extend to general directed graphs and subsets of nodes.
 - To check conditional independence of $B \perp C | A$ for subsets of nodes A , B and C . Consider all possible paths from any node in B to any node in C . Any such path is **blocked** if it includes a node such that either
 - The arrows on the path are either *tail-to-tail* or *head-to-tail* at the node, and the node is in the set A (*observed*), or
 - The arrows meet *head-to-head* at the node, and **neither the node, nor any of its descendants are in the set A (i.e., observed).**
 - If all paths are blocked, then B is d-separated from C by A .

B & C are **not** d-separated because F is not observed (tail-to-tail) and descendant of head-to-head node E , i.e., A is observed.



B & C are d-separated because F is observed (tail-to-tail) and head-to-head node E or its descendants are not observed.



Q: What if A is also Observed?

D-separation can be computed in linear time



Historical Perspective - 1

- Formalization of ideas
 - Graphical structures . . . A historical perspective from a communication perspective
 - Sewall Wright (1921) . . . Developed “*path analysis*” as a means to study statistical relationships in biological data
 - 1960’s . . . Statisticians use graphs to describe restrictions in log-linear statistical models
 - Gallagher (1963) . . . Error correcting codes as probabilistic graphs
 - Viterbi algorithm (Forney, 1973)



Historical Perspective - 2

- AI literature
 - Taxonomic hierarchies (Woods, 1975)
 - Medical diagnosis (Spiegelharter, 1990)
 - Exact algorithms for computing the joint probability distribution (Lauritzen and Spiegelharter, 1988; Pearl, 1986)
 - Learning parameters in graph-based log-linear models (Hinton and Sejnowski, 1986)
 - Bayesian networks (belief networks, causal networks or inference diagrams)
 - Approximate algorithm based on Monte Carlo methods
 - Helmholtz machines
 - Variational techniques

See books by Frey and M.I. Jordan
Also, Bishop's book and the book
by Koller and Freidman


Similar to GMM we discussed earlier



Factor Graphs & Markov Random Fields

- Factor graphs


- Suppose $P(\underline{u}, s, \underline{v}, \underline{y}) = P(\underline{u})P(s|\underline{u})P(\underline{v}|s, \underline{u})P(\underline{y}|\underline{v})$ Show Bi-partite graph



$$\{\underline{u}, s\} \perp \underline{y} | \underline{v}$$

All paths that pass from \underline{u} and s to \underline{y} pass through \underline{v}

- Markov random fields (MRF).... Undirected graphs



Markov chains in higher dimensions

Hammersley-Clifford Theorem

- Properties of MRF

- Undirected graph with nodes corresponding to variables

- $P(z_i | \mathbf{z} \setminus z_i) = P(z_i | n_i)$

z_i variable

n_i neighbors of variable z_i

\mathbf{z} = set of all variables (e.g., $\mathbf{z} = \{\underline{u}, s, \underline{v}, \underline{y}\}$)

- Local Markov Property
- Clique-based Factorization
- Global Markov Property (D-separation)

“Given its neighbors, each variable is independent of all other variables”

- Joint Distribution is given by Hammersley-Clifford Theorem (1971)

$P(\mathbf{z})$ = product of clique potentials

- o **Clique is a fully connected sub-graph that cannot remain fully connected if more variables are included**

- o **Cliques in the graphs**

$C_1 = \{\underline{u}, s, \underline{v}\} \quad C_2 = \{\underline{v}, \underline{y}\}$

$$P(\mathbf{z}) = \alpha \prod_{j=1}^{NC} \psi_j(C_j) = \alpha \exp\left\{-\sum_{j=1}^{NC} E_j(C_j)\right\}$$

NC = Number of cliques

α = Normalization Constant

Typical model of $\psi_j(C_j)$

$$\psi_j(C_j) = \exp(\{-E_j(C_j)\})$$

$E_j(C_j)$ = *energy function*

Factor Graph Example

- For the channel coding example

$$- P(\underline{u}, s, \underline{v}, \underline{y}) = \alpha \psi_1(\underline{u}, s, \underline{v}) \psi_2(\underline{v}, \underline{y})$$

when $\alpha = \frac{1}{P(\underline{v})}$

cliques

separator

$$\psi_1 = P(\underline{u}, s, \underline{v}) \quad \text{and}$$

$$\psi_2 = P(\underline{v}, \underline{y})$$

we obtain the joint distribution

Joint distribution = product of clique potentials/Product of separator potentials

Ising Model : Image-denoising

- Observed *noisy* image described by an array of binary pixel values

$$y_i \in \{-1, +1\}, i = 1, 2, \dots, p$$

Variation: $p(y_i/x_i)$ *Gaussian*

- Original (hidden) image has binary pixel values

$$x_i \in \{-1, +1\}, i = 1, 2, \dots, p$$

- Joint distribution $p(\underline{x}, \underline{y})$ is the *Boltzmann* distribution

$$p(\underline{x}, \underline{y}) = \alpha \exp\{-E(\underline{x}, \underline{y})\}$$

$$E(\underline{x}, \underline{y}) = \underbrace{\sum_{i=1}^p h_i x_i}_{\text{to bias towards } -1 \text{ or } +1} \underbrace{- \sum_{i=1}^p \sum_{j \in n_i} \beta_{ij} x_i x_j}_{\text{want energy to be small when } x_i \text{ and } x_j \text{ have same sign}} \underbrace{- \sum_{i=1}^p \eta_i x_i y_i}_{\text{want energy to be small when } x_i \text{ and } y_i \text{ have same sign}}$$

- MAP estimate via mean field (variational approximation)

$$q(\underline{x}) = \prod_{i=1}^p q_i(x_i, \mu_i); \mu_i = \text{mean value of pixel } i$$

$$\log q_i(x_i) = E_{-q_i} \{\log(p(\underline{x}, \underline{y}))\} + \text{const}$$

Mean Field Method

- Keep all pixels $j \neq i$ at their mean values

$$\log q_i(x_i) = E_{q_{-i}}[\log p(\underline{x}, \underline{y})] = x_i \left(\sum_{j \in n_i} \beta_{ij} \mu_j - h_i + \eta_i y_i \right) + \text{const}$$

$$\therefore q_i(x_i) \propto \exp(x_i \left[\left(\sum_{j \in n_i} \beta_{ij} \mu_j - h_i \right) + \eta_i y_i \right])$$


$$\Rightarrow q_i(x_i = 1) \propto \exp(a_i); q_i(x_i = -1) \propto \exp(-a_i); a_i = \left(\sum_{j \in n_i} \beta_{ij} \mu_j - h_i \right) + \eta_i y_i$$

- Update μ_i and iterate until convergence

$$\mu_i = q_i(x_i = 1) - q_i(x_i = -1) = \frac{e^{a_i} - e^{-a_i}}{e^{a_i} + e^{-a_i}} = \tanh(a_i)$$

- It is usually good to *low pass filter* the updates ($\lambda \leq 0.5$)


$$\mu_i^{t+1} = \lambda \mu_i^t + (1 - \lambda) \tanh(a_i^t)$$



Iteration 1 Iteration 3 Iteration 15

Directed Acyclic Graphs

- Bayesian Networks
 - Represented in terms of directed acyclic graphs




$$P(\underline{u}, s, \underline{v}, \underline{y}) = P(\underline{u})P(s | \underline{u})P(\underline{v} | s, \underline{u})P(\underline{y} | \underline{v})$$

- $\mathbf{z} = [z_1 \ z_2 \ \cdots \ z_N]$
 $P(z_k | a_k) \quad a_k = \text{parents of } z_k = pa(z_k)$
$$P(\mathbf{z}) = \prod_{k=1}^N P(z_k | pa(z_k)) = \prod_{k=1}^N P(z_k | a_k)$$

Chain Rule for Bayesian Networks

- Example 1



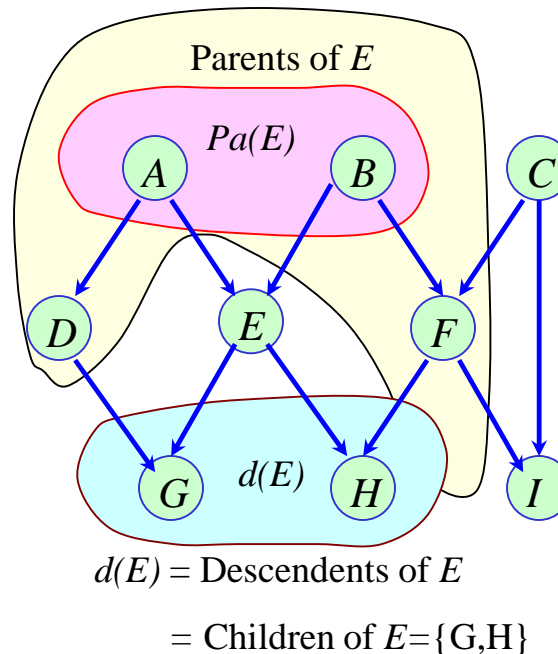
– Topological order: $(z_1 \ z_2 \ z_3 \ \cdots \ z_{13})$

$$\begin{aligned} P(\mathbf{z}) = & P(z_1).P(z_2 | z_1).P(z_3).P(z_4 | z_1, z_2).P(z_5 | z_2, z_3). \\ & P(z_6).P(z_7 | z_3).P(z_8 | z_4, z_5).P(z_9 | z_5, z_6).P(z_{10}). \\ & P(z_{11} | z_8).P(z_{12} | z_8, z_9).P(z_{13} | z_9, z_{10}) \end{aligned}$$

Markov Blanket

- Example 2:

Markov Blanket of a node E is denoted by ∂E . ∂E = its parents, its children, and its children's other parents.



$\{A, B\}$ = Parents of E

$\{D, F\}$ = its children's other parents

$$P(E | \partial E, C, I) = P(E | \partial E)$$

$$\partial E = \{A, B, D, F, G, H\}$$

$$P(A, B, C, D, E, F, G, H, I) = P(A)P(B)(C)P(D|A)P(E|A, B)P(F|B, C)$$

$$P(G|D, E)P(H|E, F)P(I|C, F)$$

For an undirected graph (Markov Random Field), Markov Blanket of a node is the set of its neighboring nodes



Topological Ordering

- Can arrange nodes in topological order
 - For each node x all of its parents $pa(x)$ precede it in the ordering
- Topological orders are not unique
 - Order 1: $\{A, B, C, D, E, F, G, H, I\}$
 - Order 2: $\{B, A, E, D, G, C, F, I, H\}$



Constructing Topological Ordering

- Algorithms for finding topological ordering
 - Algorithm 1:
 - Start with the graph and an empty list
 - Successively delete from the graph any node which does not have any parents, and add it to the end of the list
 - Stop when no node has parent nodes
 - Algorithm 2:
 - Start with the graph and an empty list
 - Successively delete from the graph nodes which have no children and add them to the beginning of the list
 - Stop when no node has child node

Inference via Variable Elimination

- Consider the out-tree structure

$$P(a, b, c, d, e, f, g)$$

$$= P(a)P(b | a)P(c | a)P(d | b)P(e | b)P(f | c)P(g | c)$$

- Suppose you observe e and f
- Want to compute $P(d | e, f)$


$$P(d | e, f) = \sum_{a, b, c, g} P(d, a, b, c, g | e, f) \longrightarrow$$

Too much computation by brute force

$$= \sum_b P(d | b) \sum_a P(b | a, e) \sum_c P(a | c) \sum_g P(c, g | f)$$


$$= \sum_b P(d | b) \sum_a P(b | a, e) \sum_c P(a | c) P(c | f)$$

- Ordered summation: c, a, b
- What if the graph is a general directed acyclic graph (DAG)?



Variable Elimination for DAGs


- Consider the Directed Acyclic Graph (DAG).
Want to find the marginal probability $P(g)$
- Steps involve sums and products
 - Compute the product $P(a, b, d) = P(a)P(b)P(d | a, b)$
 - Sum over b to get $P(a, d) = \sum_b P(a, b, d)$
 - Multiply $P(a, d)$ by $P(c | a)$ to get $P(a, c, d) = P(c | a)P(a, d)$
 - Sum $P(a, c, d)$ over a to get $P(c, d) = \sum_a P(a, c, d)$
 - Multiply $P(c, d)$ by $P(e | c, d)$ to obtain $P(c, d, e) = P(e | c, d)P(c, d)$
 - Sum over c and d to get $P(e) = \sum_c \sum_d P(c, d, e)$
 - Multiply $P(e)$ by $P(g | e)$ to get $P(e, g)$
 - Sum $P(e, g)$ over e to get $P(g) = \sum_e P(e, g)$
- Complexity is exponential in the size of factors and optimal ordering of computations is NP-hard
- Is there a formal (and nicer) way to do inference in Bayesian networks? For trees, there is a nice *sum-product algorithm* as in HMMs. For general DAGs, *Junction tree algorithm*.





Sum-Product Algorithm using Factor Graphs

- Trees: single undirected path between each pair of nodes



Joint probability distribution as a product of factors

$$p(\underline{x}) = \prod_s f_s(\underline{x}_s); s = \text{factor}; \underline{x}_s = \text{subset of variables in factor } s$$

$$\text{Also, } p(\underline{x}) = \prod_{s \in ne(x)} F_s(x, \underline{X}_s); \underline{X}_s = \text{set of all variables in the subtree connected to } x \text{ via } s$$

$$\text{Example: } p(\underline{x}) = f_a(x_1, x_2) f_b(x_2, x_3) f_c(x_2, x_4)$$

$$\text{For } x_1: p(\underline{x}) = F_a(x_1, \underbrace{x_2, x_3, x_4}_{\underline{X}_a})$$


$$= f_a(x_1, x_2) f_b(x_2, x_3) f_c(x_2, x_4)$$

$$\text{For } x_2: p(\underline{x}) = \underbrace{f_a(x_1, x_2)}_{F_a(x_2, X_a)} \underbrace{f_b(x_2, x_3)}_{F_b(x_2, X_b)} \underbrace{f_c(x_2, x_4)}_{F_c(x_2, X_c)}$$

$$X_a = x_1, X_b = x_3, X_c = x_4$$

so, $|ne(x_i)| = \text{number of terms involving } x_i \text{ in } p(\underline{x})$

Factor graphs are bi-partite graphs





Message Passing between Variables and Factors: Example

- **Factors → Variables Messages (SUM)**

messages sent by a factor node to a variable node involves multiplying all the incoming messages (except variable node x) with the factor and summing over all the variables except x

$$\mu_{f_a \rightarrow x_1}(x_1) = \sum_{x_2} f_a(x_1, x_2) \mu_{x_2 \rightarrow f_a}(x_2); \mu_{f_a \rightarrow x_2}(x_2) = \sum_{x_1} f_a(x_1, x_2) \mu_{x_1 \rightarrow f_a}(x_1)$$

$$\mu_{f_b \rightarrow x_2}(x_2) = \sum_{x_3} f_b(x_2, x_3) \mu_{x_3 \rightarrow f_b}(x_3); \mu_{f_b \rightarrow x_3}(x_3) = \sum_{x_2} f_b(x_2, x_3) \mu_{x_2 \rightarrow f_b}(x_2)$$

$$\mu_{f_c \rightarrow x_2}(x_2) = \sum_{x_4} f_c(x_2, x_4) \mu_{x_4 \rightarrow f_c}(x_4); \mu_{f_c \rightarrow x_4}(x_4) = \sum_{x_2} f_c(x_2, x_4) \mu_{x_2 \rightarrow f_c}(x_2)$$

- **Variables → Factors Messages (PRODUCT)**


Message sent by a variable node to a factor node is the product of all the incoming messages along all of the other links (factors)

$$\mu_{x_1 \rightarrow f_a}(x_1) = 1; \mu_{x_3 \rightarrow f_b}(x_3) = 1; \mu_{x_4 \rightarrow f_c}(x_4) = 1$$

$$\mu_{x_2 \rightarrow f_a}(x_2) = \mu_{f_b \rightarrow x_2}(x_2) \mu_{f_c \rightarrow x_2}(x_2);$$

$$\mu_{x_2 \rightarrow f_b}(x_2) = \mu_{f_a \rightarrow x_2}(x_2) \mu_{f_c \rightarrow x_2}(x_2);$$

$$\mu_{x_2 \rightarrow f_c}(x_2) = \mu_{f_a \rightarrow x_2}(x_2) \mu_{f_b \rightarrow x_2}(x_2)$$





Message Passing between Factors and Variables

- Use the node for which you want to compute marginal probability as the root

- **Factors \rightarrow Variable Messages**

Marginal probability of a variable x

$$p(x) = \prod_{s \in ne(x)} \sum_{\underline{X}_s} F_s(x, \underline{X}_s) = \prod_{s \in ne(x)} \mu_{f_s \rightarrow x}(x)$$

Recursively,



$$F_s(x, \underline{X}_s) = f_s(x, \underbrace{x_1, \dots, x_M}_{\underline{X}_s}) G_1(x_1, \underline{X}_{s1}) \dots G_M(x_M, \underline{X}_{sM})$$

$$\mu_{f_s \rightarrow x}(x) = \sum_{x_1} \dots \sum_{x_M} f_s(x, x_1, \dots, x_M) \underbrace{\prod_{m \in ne(f_s) \setminus x} \sum_{\underline{X}_{sm}} G_m(x_m, \underline{X}_{sm})}_{\mu_{x_m \rightarrow f_s}(x_m)}$$

$$= \sum_{x_1} \dots \sum_{x_M} f_s(x, x_1, \dots, x_M) \prod_{m \in ne(f_s) \setminus x} \mu_{x_m \rightarrow f_s}(x_m)$$

If factor s is a leaf node with only variable x ,

set $\mu_{f_s \rightarrow x}(x) = f_s(x)$



Messages passed along a link are always a function of the variable it is connected to





Message Passing between Variables and Factors

- Variables → Factors Messages**

$$\begin{aligned}
 \mu_{x_m \rightarrow f_s}(x_m) &= \sum_{\underline{X}_{sM}} G(x_M, \underline{X}_{sM}) \\
 &= \sum_{\underline{X}_{sM}} \prod_{l \in ne(x_m) \setminus f_s} F_l(x_m, \underline{X}_{ml}) \\
 &= \prod_{l \in ne(x_m) \setminus f_s} \sum_{\underline{X}_{ml}} F_l(x_m, \underline{X}_{ml}) \\
 &= \prod_{l \in ne(x_m) \setminus f_s} \mu_{f_l \rightarrow x_m}(x_m)
 \end{aligned}$$

If x_m is a leaf node, $\mu_{x_m \rightarrow f_s}(x_m) = 1$

- Message sent by a variable node to a factor node is the product of all the incoming messages *along all of the other links* (factors)
- On the other hand, messages sent by a factor node to a variable node involves multiplying all the incoming messages (except variable node x) with the factor and summing over all the variables except x





Computing All Marginal Probabilities

- Select an arbitrary node as the root and propagate messages from the leaves to the root as in the sum-product algorithm for a single root node
- Send messages from the root all the way back to the leaves
- Now calculate the marginal probability at each variable and factor node via

$$p(x) = \prod_{s \in ne(x)} \mu_{f_s \rightarrow x}(x); \quad p(\underline{x}_s) = f_s(\underline{x}_s) \prod_{x_i \in f_s} \mu_{x_i \rightarrow f_s}$$

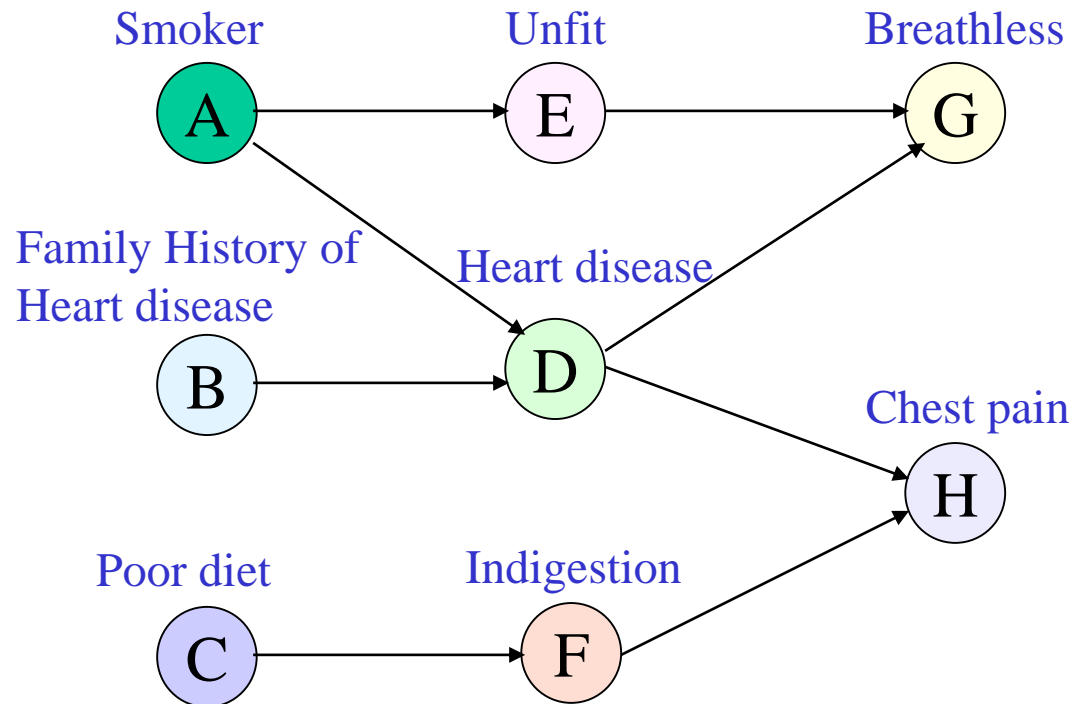
- Can eliminate messages from variable nodes to factors via

$$\begin{aligned} \mu_{f_s \rightarrow x}(x) &= \sum_{x_1} \dots \sum_{x_M} f_s(x, x_1, \dots, x_M) \prod_{m \in ne(f_s) \setminus x} \mu_{x_m \rightarrow f_s}(x_m) \\ &= \sum_{x_1} \dots \sum_{x_M} f_s(x, x_1, \dots, x_M) \prod_{m \in ne(f_s) \setminus x} \left(\prod_{l \in ne(x_m) \setminus f_s} \mu_{f_l \rightarrow x_m}(x_m) \right) \end{aligned}$$

- MAP problem is called the Max-sum algorithm (Viterbi for Trees)
- For DAGs and MRFs, multiple paths may exist. If you use sum-product the usual way, it is called *Loopy Belief Propagation* and it works OK!

General DAGs

- Constructing the Inference Engine:
 - Consider an artificial medical diagnosis problem



$$P(ABCDEFGH) = P(A)P(B)P(C)P(E | A)P(D | A, B)P(F | C) \\ P(G | E, D)P(H | D, F)$$



Inference Problem in DAGs

- Typically, we are interested in computing the marginal distributions conditioned on some observation of one or more variables
- Example: what is the probability of Heart disease Given that the patient is a smoker, is Breathless and has chest pain?

$$P(D = T \mid A = G = H = T) \qquad T \Rightarrow \text{TRUE}$$

How to compute inference probabilities efficiently?




Key Steps in Inference for DAGs

- Key steps in exact Bayesian inference
 1. Add undirected edges to all co-parents which are not currently joined (a process called *marrying parents*)
 2. Drop all directions in the graph obtained from stage 1. The result is the so-called *moral graph*.
 3. Triangulate the *moral graph*, that is, add sufficient additional undirected links between nodes such that *there are no cycles (i.e., closed paths) of length 4 or more distinct nodes without a short-cut*.
 4. Identify the *cliques* of this triangulated graph
 5. Join the cliques together to form the *junction tree*
 6. Perform inference on the junction tree (*message passing*)



Marrying Parents


- Step 1: Marrying parents





Obtaining the Moral Graph

- Step 2: Moral graph

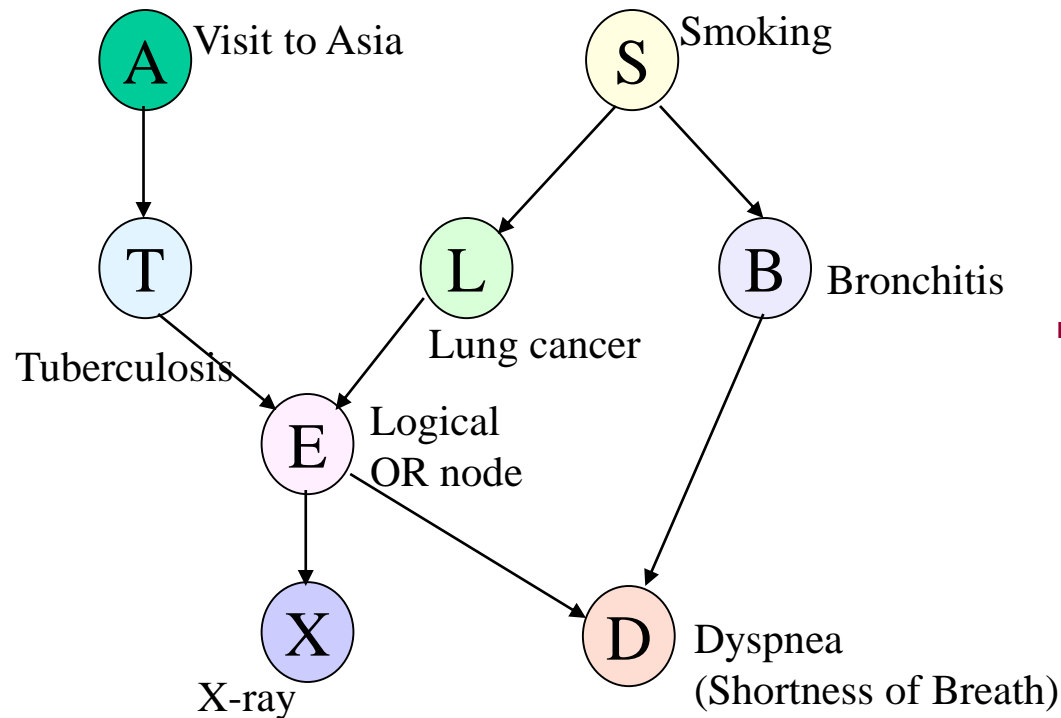


Triangulating the Moral Graph


- Step 3: Triangulate the moral graph

It is already triangulated.

Example where triangulation is needed:



Original graph for Asia Problem




Moral graph



Triangulated Graph for the Asia Problem

$$P(ASTLBEXD) = P(A)P(S)P(T | A)P(L | S)P(B | S) \\ P(E | LT)P(X | E)P(D | B, E)$$



Triangulated graph



Cliques of the Triangulated Graph

Step 4: Cliques of the triangulated graph

Clique: a fully connected (complete) maximal subgraph

Medical Diagnosis Problem:


C_1 : ABD

C_2 : ADE

C_3 : DEG

C_4 : DFH

C_5 : CF



Asia Problem:

C_1 : AT


C_2 : TLE

C_3 : BLE

C_4 : SBL

C_5 : DBE

C_6 : XE






Constructing the Junction Tree

- Step 5: Make the junction tree

Key property: **running intersection property** \Rightarrow If a variable x is contained in two cliques, then it is contained in every clique on the path connecting the two cliques.

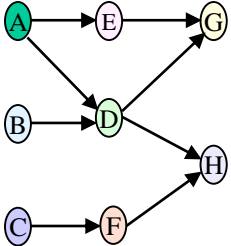
The edge joining two cliques is called a **separator**.





Joint Distribution in terms of Cliques and Separators

KEY: $P(ABCDEFGH) = \frac{\prod_{i=1}^5 P(C_i)}{\prod_{i=2}^5 P(S_i)}$ “Marginal representation”



$$= \frac{P(ABD)P(ADE)P(DEG)P(DFH)P(CF)}{P(AD)P(DE)P(D)P(F)}$$

Recall that

$$P(ABCDEFGH) = P(A)P(B)P(C)P(D | AB)P(E | A)P(F | C) \\ P(G | DE)P(H | DF)$$

Note that

$$P(C_1) = P(ABD) = P(D | AB)P(A)P(B)$$

$$P(C_2) = P(ADE) = P(E | AD)P(AD) = P(E | A) \cdot P(S_2)$$

$$P(C_3) = P(G | DE) \cdot P(DE) = P(G | DE) \cdot P(S_3)$$



Joint Distribution in terms of Clique Potentials

$$P(C_4) = P(H \mid FD) \cdot P(S_4) \cdot P(S_5)$$


$$P(C_5) = P(C \mid F) \cdot P(F) = P(F \mid C) \cdot P(C)$$

So, marginal representation does indeed provide the joint distribution. In fact,

$$\text{Separator } S_i = C_i \cap \{C_1 \cup C_2 \cup \dots \cup C_{i-1}\}$$

$$\text{Let } R_i = C_i \setminus S_i \Rightarrow C_i - S_i$$

$$\begin{aligned} S_3 &= \{DEG\} \cap \{\{ABD\} \cup \{ADE\}\} \\ &= \{DEG\} \cap \{ABDE\} = DE \\ R_3 &= G \end{aligned}$$



$$P(ABCDEFGH) = P(C_1) \prod_{i=2}^5 P(C_i \mid S_i)$$

$$= P(C_1) \prod_{i=2}^5 P(R_i \mid S_i) = \prod_{i=1}^5 P(R_i \mid S_i); S_1 = \phi$$

$$= P(ABD)P(E \mid AD)P(G \mid DE)P(HF \mid D)P(C \mid F)$$


$$= \prod_{i=1}^5 \psi(C_i)$$

This is called a *potential representation* of joint distribution



Non-uniqueness of Junction Trees

Junction tree is not unique:



$$P(ABCDEFGH) = P(ABD)P(E | AD)P(G | DE)P(HF | D)P(C | F)$$

$$= \prod_{i=1}^5 \psi(C_i)$$

$$= P(A)P(B)P(D | AB) \frac{P(C)P(F | C)}{P(F)} P(E | A)P(G | DE)P(H | DF)P(F)$$

Note that ψ can be any function of cliques with a suitable normalization at the end


$$\text{Example : } \bar{\psi}(C_i) = \psi(C_i) \sum_{R_i} \psi(C_i) = P(R_i | S_i)P(S_i) = P(R_i, S_i)$$

$$P(R_i | S_i) = \frac{P(R_i, S_i)}{P(S_i)} = \frac{\bar{\psi}(C_i)}{\sum_{R_i} \bar{\psi}(C_i)} \Rightarrow \text{can work with } \bar{\psi}(C_i) \text{ directly}$$




Junction Tree for the Asia Problem

One junction tree for the Asia Problem:



Junction graph



Triangulated graph



Cliques, Separators & Potentials

$$P(ASTLBEXD) = P(A)P(S)P(T | A)P(L | S)P(B | S)$$

$$P(E | LT)P(X | E)P(D | BE)$$

Chain Rule

$$\text{Cliques \& Separators} = \frac{P(AT)P(TLE)P(BLE)P(SBL)P(DBE)P(XE)}{P(T)P(LE)P(BL)P(BE)P(E)}$$

Clique Potentials/
Factors

$$= P(AT)P(LE | T)P(B | LE)P(S | BL)$$

$$P(D | BE)P(X | E)$$

$$= \prod_{i=1}^6 \psi(C_i)$$

Although it looks strange, it does work!!!

- Step 6: Inference on the junction tree: sum-product algorithm on cliques

for each node in the junction tree, store

clique, $R_i (= C_i \setminus S_i)$, S_i and $\psi(\text{clique})$



Other Methods for Inference

- Bayesian Inference
 - The junction tree approach becomes intractable for dense graphs
 - Alternate Approaches
 - Probabilistic logic sampling on
 - » DAGs
 - » Junction tree
 - **Gibbs sampling**
 - Boltzmann Machines
 - » Gibbs sampling
 - » Mean Field Approximation
 - Lagrangian Relaxation (Variational approximation)
 - Expectation Propagation
- Learning BN Parameters and Structure from Data



What is a Gibbs Sampler?

- It is a Markov Chain Monte Carlo method (recall particle filter)
 - Updates one variable at a time
 - Samples from a conditional distribution of a variable when other variables are fixed
 - Ideally suited for Bayesian networks
- Suppose you want to sample from a distribution of p variables $p(x_1, x_2, \dots, x_p)$
 - Initialize $\{x_i^0\}_{i=1}^p$
 - For $t = 1, 2, \dots, T$
 - Sample $x_1^{(t+1)} \sim p(x_1 | x_2^t, x_3^t, \dots, x_p^t)$
 - Sample $x_2^{(t+1)} \sim p(x_2 | x_1^{(t+1)}, x_3^t, \dots, x_p^t)$
 -
 - Sample $x_i^{(t+1)} \sim p(x_i | x_1^{(t+1)}, \dots, x_{i-1}^{(t+1)}, x_{i+1}^{(t)}, \dots, x_p^t)$
 -
 - Sample $x_p^{(t+1)} \sim p(x_p | x_1^{(t+1)}, x_2^{(t+1)}, \dots, x_{p-1}^{(t+1)})$

- Need a burn-in period
- Subsample to minimize correlations



Summary

- Graphical Models
- Bayesian Inference in Graphical Models
- Forward-Backwards Methods of Inference
- Simulation-based Methods