

Non-parametric Learning of Weights

- Perceptron (linear) and LVQ (non linear)
- Relaxation Procedures
- Widrow – Hoff (LMS, Adaline)
- Ho-Kashyap Procedure
- Incremental Gauss – Newton = RLS
- Fisher’s linear discriminant
- Support Vector Machines

One can trace the history of Neural networks to the work of Santiago Ramon y Cajal, who discovered that the basic building element of the brain is the neuron. The brain comprises approximately 60-100 billion neurons! Each neuron is connected with other neurons via elementary structural and functional units/links, known as synapses. It is estimated that there are 50-100 trillion synapses. These links mediate information between connected neurons.

The most common type of synapses are the chemical ones, which convert electrical pulses, produced by a neuron, to a chemical signal and then back to an electrical one. Depending on the input pulse(s), a synapse is either activated or inhibited. Via these links, each neuron is connected to other neurons and this happens in a hierarchical way, in a layer-wise fashion.

In 1943, Warren McCulloch and Walter Pitts, developed a computational model for the basic neuron linking neurophysiology with mathematical logic. They showed that given a sufficient number of neurons and adjusting appropriately the synaptic links, each one represented by a weight, one can compute, in principle, any computable function. As a matter of fact, it is generally accepted that this is the paper that gave birth to the fields of neural networks and artificial intelligence.

Frank Rosenblatt borrowed the idea of a neuron model, as suggested by McCulloch and Pitts, and proposed a true learning machine, which learns from a set of training data. In its most basic version, he used a single neuron and adopted a rule that can learn to separate data, which belong to two linearly separable classes. That is, he built a Pattern Recognition system. He called the basic neuron a perceptron and developed a rule/algorithm, the perceptron algorithm, which we used in HW 1 and review it briefly.

Discuss a neuron and a relay ... idealization of $\sigma(c, \underline{x}) = \frac{1}{1 + \exp(-c \underline{w}^T \underline{x})}$... as $c \rightarrow \infty$, $step = g(\underline{x})$

$$\frac{d\sigma(c, \underline{x})}{dx} = c\sigma(c, \underline{x})(1 - \sigma(c, \underline{x}))$$

when use mse criterion or cross – entropy

$$gradient = - \sum_{n=1}^N e_n \underline{x}_n = - \sum_{n=1}^N (z_n - \sigma(c, \underline{x}_n)) \underline{x}_n \sigma(c, \underline{x}_n) (1 - \sigma(c, \underline{x}_n))$$

Formalization:

Select class 1 if $\underline{w}^T \underline{x} > 0$; $\underline{x} = (-1 \ x_1 \ x_2 \ \cdots \ x_p)$

Look at it geometrically by recalling that the inner product is related to cosine of the angle between vectors.

$$\underline{w}^T \underline{x} > 0 \text{ if } \underline{x} \text{ belongs to class 1} \quad \Rightarrow \theta < 90^\circ \text{ (acute angle)}$$

$$\underline{w}^T \underline{x} < 0 \text{ if } \underline{x} \text{ belongs to class 2} \quad \Rightarrow \theta > 90^\circ \text{ (obtuse angle)}$$

$$\text{since } \underline{w}^T \underline{x} = \|\underline{w}\| \|\underline{x}\| \cos \theta$$

You can setup an optimization problem to minimize the number of errors or a measure of errors (e.g., how far away from 0).

$$\sum_{n=1}^N [z_n \max(0, b_n - \underline{w}^T \underline{x}_n) + (1 - z_n) \max(0, \underline{w}^T \underline{x}_n + b_n)]$$

$$b_n = 0 \Rightarrow \sum_{n=1}^N [z_n \max(0, -\underline{w}^T \underline{x}_n) + (1 - z_n) \max(0, \underline{w}^T \underline{x}_n)]$$

$$\text{Error metric: } \Rightarrow z_n = 1 \& \underline{w}^T \underline{x}_n > 0 \Rightarrow \text{no cost} \& z_n = 1 \& \underline{w}^T \underline{x}_n < 0 \Rightarrow \text{gradient} - \underline{x}_n$$

$$z_n = 0 \& \underline{w}^T \underline{x}_n < 0 \Rightarrow \text{no cost} \& z_n = 0 \& \underline{w}^T \underline{x}_n > 0 \Rightarrow \text{gradient } \underline{x}_n$$

same conclusions apply with nonzero b_n !

Perceptron: Supervisory Learning (Reinforcement Learning) using **incremental gradient** or **stochastic gradient**... Update \underline{w} only if you make mistakes

Do until errors stabilize

Do $n = 1 : N$

if $e_n \neq 0$

$$\underline{w} \leftarrow \underline{w} + \eta e_n \underline{x}_n$$

end if

end

end

$$\text{if } z_n = 1 \text{ and } g(\underline{x}_n) = 0, e_n = 1 \Rightarrow \underline{w}^T \underline{x}_n \leftarrow \underline{w}^T \underline{x}_n + \eta \underline{x}_n^T \underline{x}_n \quad \uparrow \text{ as it should!}$$

$$\text{if } z_n = 0 \text{ and } g(\underline{x}_n) = 1, e_n = -1 \Rightarrow \underline{w}^T \underline{x}_n \leftarrow \underline{w}^T \underline{x}_n - \eta \underline{x}_n^T \underline{x}_n \quad \downarrow \text{ as it should!}$$

Perceptron learning rule (Incremental gradient, Stochastic gradient)

Alternate: Compute mini-batch gradient

Do until errors stabilize
Do $n = 1 : N$ in steps of m

$$\underline{w} \leftarrow \underline{w} + \eta \sum_{j=n}^{n+m-1} e_n \underline{x}_n$$

end if
end
end

Perceptron Convergence Theorem:

For simplicity, scale class 0 samples as $\underline{x}^i = -\underline{x}^i$

Go through the two proofs

Optimal Learning rates or relaxation method

Extension to Multiple Classes

Key properties of Perceptron:

- It employs distributed decision-based credit assignment
- Update weights only when misclassification occurs
- Distributed and localized: reinforce correct class, penalize wrong decision class
- Update depends on incremental/stochastic gradient

How many random patterns a Perceptron with p inputs can learn reliably in a 2 class case? $2p$

$P(N, p) = \begin{cases} 1 & N \leq p+1 \\ \frac{2}{2^N} \sum_{i=0}^p \binom{N-1}{i} & N > p+1 \\ \approx \Phi\left(\frac{2p-N}{\sqrt{N}}\right) & \text{for large } N \end{cases}$	$N=4; p=2$ $\Rightarrow P(N, p) = \frac{1}{8}(1+3+3)$ $= \frac{7}{8}$ <i>XOR Pattern cannot be correctly classified by a Perceptron</i>
-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------	--------------------------------------------------------------------------------------------------------------------------------------------------

Back Propagation:

$$L(x) = L(f(h(g(x))))$$

$$\frac{\partial L}{\partial x} = \frac{\partial L}{\partial f} \frac{\partial f}{\partial h} \frac{\partial h}{\partial g} \frac{\partial g}{\partial x} = \lambda_g \frac{\partial g}{\partial x}$$

$$\lambda_f = \frac{\partial L}{\partial f}; \lambda_h = \frac{\partial L}{\partial h} = \lambda_f \frac{\partial f}{\partial h};$$

$$\lambda_g = \frac{\partial L}{\partial g} = \lambda_h \frac{\partial h}{\partial g}$$

$$x \rightarrow g \rightarrow h \rightarrow f \rightarrow L$$

$$\frac{\partial L}{\partial x} = \lambda_g \frac{\partial g}{\partial x} \leftarrow \lambda_g = \frac{\partial L}{\partial g} = \lambda_h \frac{\partial h}{\partial g} \leftarrow \lambda_h = \frac{\partial L}{\partial h} = \lambda_f \frac{\partial f}{\partial h} \leftarrow \lambda_f = \frac{\partial L}{\partial f}$$

Decision Trees:

- Developed since 1960s and popularized by Brieman and Quinlan
 - \Rightarrow Tree Algorithms: ID3, C4.5, C5.0 and CART
- Three uses: Data description (compression, rules), classification and generalization
- Popular in statistics, pattern recognition, decision theory, signal processing and machine learning
- Have been used in industrial applications, particularly in diagnosis and quality control. For example, look at my papers on sequential fault diagnosis since 1990. Primarily in IEEE T-SMC.
- Why? Invariant to scaling; Can handle large datasets; Easily ignore redundant variables; Handle missing variables through surrogate splits; Easy to interpret and explain
- Nearly half of the data mining competitions are won by using some variants of tree ensemble methods: Boosting, Random Forests, Bagging
- Surveys: Srirama K. Murthy: "Automatic Construction of Decision Trees From Data: A Multi-disciplinary Survey," pp 1-49. Kulwer Academic Publishers. Data Mining and Knowledge Discovery 2 (4): 345-389 (1998)

S. B. Kotsiantis, "Decision trees: a Recent Overview," Artificial Intelligence Review, 39:261–283, 2013.

Tests: single attribute tests; hyperplane tests

Continuous tests: Ranges

Select t such that mutual information $IG(z|x, t)$ is maximum

$$IG(z|x, t) = H(z) - H(z|x, t)$$

$$H(z) = -\sum_{j=1}^C P(z=j) \log_2 P(z=j)$$

$$H(z|x, t) = P(x < t)H(z|x < t) + P(x \geq t)H(z|x \geq t)$$

$$z \in \{1, 2, \dots, C\}; \underline{x} \in R^n; p(\underline{x}) = \sum_{i=1}^C P(z=i) p(\underline{x}|z=i) = \sum_{i=1}^C \pi_i N(\underline{x}; \mu_i, \sigma_i^2)$$

Recall for a binary test on x with a threshold t : $IG(z|x, t) = H(z) - H(z|x, t)$

$$H(z) = -\sum_{j=1}^C P(z=j) \log_2 P(z=j) = -\sum_{j=1}^C \pi_j \log_2 \pi_j$$

$$H(z|x, t) = P(x < t)H(z|x < t) + P(x \geq t)H(z|x \geq t) = \left(\sum_{i=1}^C \pi_i \Phi\left(\frac{t - \mu_i}{\sigma_i}\right) \right) H(z|x < t) + \left(\sum_{i=1}^C \pi_i [1 - \Phi\left(\frac{t - \mu_i}{\sigma_i}\right)] \right) H(z|x \geq t)$$

where $\Phi(c) = \frac{1}{\sqrt{2\pi}} \int_{-\infty}^c e^{-u^2/2} du = CDF$ of a standard Normal (Gaussian) distribution.

$$\text{Note that by Bayes rule, } P(z=j|x < t) = \frac{P(x < t|z=j)P(z=j)}{P(x < t)} = \frac{\Phi\left(\frac{t - \mu_j}{\sigma_j}\right)\pi_j}{\sum_{i=1}^C \pi_i \Phi\left(\frac{t - \mu_i}{\sigma_i}\right)}$$

$$\text{So, } H(z|x < t) = -\sum_{j=1}^C \left(\frac{\Phi\left(\frac{t - \mu_j}{\sigma_j}\right)\pi_j}{\sum_{i=1}^C \pi_i \Phi\left(\frac{t - \mu_i}{\sigma_i}\right)} \right) \log_2 \left(\frac{\Phi\left(\frac{t - \mu_j}{\sigma_j}\right)\pi_j}{\sum_{i=1}^C \pi_i \Phi\left(\frac{t - \mu_i}{\sigma_i}\right)} \right); H(z|x \geq t) = -\sum_{j=1}^C \left(\frac{[1 - \Phi\left(\frac{t - \mu_j}{\sigma_j}\right)]\pi_j}{\sum_{i=1}^C \pi_i [1 - \Phi\left(\frac{t - \mu_i}{\sigma_i}\right)]} \right) \log_2 \left(\frac{[1 - \Phi\left(\frac{t - \mu_j}{\sigma_j}\right)]\pi_j}{\sum_{i=1}^C \pi_i [1 - \Phi\left(\frac{t - \mu_i}{\sigma_i}\right)]} \right)$$

$$\text{So, } IG(z|x, t) = -\sum_{j=1}^C [\pi_j \log_2 \pi_j - \Phi\left(\frac{t - \mu_j}{\sigma_j}\right)\pi_j \log_2 \left(\frac{\Phi\left(\frac{t - \mu_j}{\sigma_j}\right)\pi_j}{\sum_{i=1}^C \pi_i \Phi\left(\frac{t - \mu_i}{\sigma_i}\right)} \right) - [1 - \Phi\left(\frac{t - \mu_j}{\sigma_j}\right)]\pi_j \log_2 \left(\frac{[1 - \Phi\left(\frac{t - \mu_j}{\sigma_j}\right)]\pi_j}{\sum_{i=1}^C \pi_i [1 - \Phi\left(\frac{t - \mu_i}{\sigma_i}\right)]} \right)]$$

Discuss Entropy

Gini Index

Mutual Information

JMI

Decision Trees: Low Bias and High Variance

Cost-complexity pruning

CART

Bagging and Random Forest

Boosting: AdaBoost and Gradient Boosting

Missing Value Problem

ECC

Examples

